

# U-Net – Deep Learning for Cell Counting, Detection, and Morphometry

Thorsten Falk<sup>1,2,3,†</sup>, Dominic Mai<sup>1,2,4,†</sup>, Robert Bensch<sup>1,2,†</sup>, Özgün Çiçek<sup>1</sup>, Ahmed Abdulkadir<sup>1,5</sup>, Yassine Marrakchi<sup>1,2,3</sup>, Anton Böhm<sup>1</sup>, Jan Deubner<sup>6,7</sup>, Zoe Jäckel<sup>6,7</sup>, Katharina Seiwald<sup>6</sup>, Alexander Dovzhenko<sup>8</sup>, Olaf Tietz<sup>8</sup>, Cristina Dal Bosco<sup>8</sup>, Sean Walsh<sup>8</sup>, Deniz Saltukoglu<sup>2,9,10,11</sup>, Tuan Leng Tay<sup>7,12,13</sup>, Marco Prinz<sup>2,3,12</sup>, Klaus Palme<sup>2,8</sup>, Matias Simons<sup>2,9,10,14</sup>, Ilka Diester<sup>6,7,15</sup>, Thomas Brox<sup>1,2,3,7</sup>, and Olaf Ronneberger<sup>1,2,\*</sup>

<sup>1</sup>Department of Computer Science, Albert-Ludwigs-University, Freiburg, Germany

<sup>2</sup>BIOSS Centre for Biological Signalling Studies, Freiburg, Germany

<sup>3</sup>CIBSS Centre for Integrative Biological Signalling Studies, Albert-Ludwigs-University, Freiburg, Germany

<sup>4</sup>Life Imaging Center; Center for Biological Systems Analysis, Albert-Ludwigs-University, Freiburg, Germany

<sup>5</sup>University Hospital of Old Age Psychiatry and Psychotherapy, University of Bern, Bern, Switzerland

<sup>6</sup>Optophysiology Lab, Institute of Biology III, Albert-Ludwigs-University, Freiburg, Germany

<sup>7</sup>BrainLinks-BrainTools, Albert-Ludwigs-University, Freiburg, Germany

<sup>8</sup>Institute of Biology II, Albert-Ludwigs-University, Freiburg, Germany

<sup>9</sup>Center for Biological Systems Analysis, Albert-Ludwigs-University, Freiburg, Germany

<sup>10</sup>Renal Division, University Medical Centre, Freiburg, Germany

<sup>11</sup>Spemann Graduate School of Biology and Medicine (SGBM), Albert-Ludwigs-University, Freiburg, Germany

<sup>12</sup>Institute of Neuropathology, University Medical Centre, Freiburg, Germany

<sup>13</sup>Institute of Biology I, Albert-Ludwigs-University, Freiburg, Germany

<sup>14</sup>Paris Descartes University-Sorbonne Paris Cité, Imagine Institute, Paris, France

<sup>15</sup>Bernstein Center Freiburg, Albert-Ludwigs-University, Freiburg, Germany

<sup>†</sup>Equal Contribution

<sup>\*</sup>Corresponding author

U-Net is a generic deep-learning solution for frequently occurring quantification tasks like cell detection and shape measurements in biomedical image data. We present an ImageJ plugin that enables non-machine learning experts to analyze their data with U-Net either on the local computer or a remote server/cloud service. The plugin comes with pre-trained models for single cell segmentation and allows adaptation of U-Net to new tasks based on few annotated samples.

The advancement of microscopy and sample preparation techniques leaves researchers with large amounts of image data. The stacks of data promise additional insights, more precise analysis, and more rigorous statistics, were there not the hurdles of quantification. Images must be converted first into numbers before they are accessible for statistical analysis. Often this requires counting thousands of cells with a certain marker or drawing the outlines of cells to quantify their shape or the strength of a reporter. Such work is not much liked in the lab and, consequently, it is often avoided. In neuroscientific studies applying optogenetic tools, for example, it is often requested to quantify the number of opsin expressing neurons or the localization of newly developed opsins within the cells. However, because of the effort, most studies are published without this information.

Is such quantification not a job that computers can do? Indeed it is. For decades, computer scientists have developed specialized software that can lift the burden of quantification from researchers in the life sciences. However, each lab produces different data and focuses on different aspects of the data for their research question at hand. Thus, a new software must be built for each case. Deep learning could change this situation. It learns the features relevant for the task from data rather than being hard-coded. Therefore, software engineers need not set up a specialized software for a certain quantification task. Instead, a generic software package can learn to adapt to the task autonomously from appropriate data; data that researchers in the life sciences can provide by themselves.

Learning-based approaches caught the interest of the biomedical community already years ago. Popular solutions like ilastik<sup>1a</sup> or the trainable WEKA segmentation toolkit<sup>2b</sup> allow training of segmentation pipelines using generic hand-tailored image features. More recently, attention has shifted towards deep learning which automatically extracts optimal features for the actual image analysis task, lifting the need for feature design by experts in computer science<sup>3–5</sup>. However, a wide-spread use for quantification in the life sciences has been impaired by the lack of generic, easy-to-use software packages. While the packages Aivia<sup>c</sup> and Cell Profiler<sup>6d</sup> already use deep learning models, they do not allow training on new data. This restricts their application domain to a small range of datasets. In the scope of image restauration, CSBDeep<sup>7</sup> provides an ImgLib2<sup>e</sup>-based plugin with models for specific imaging modalities and biological specimen. It allows to integrate externally trained new restoration models. Another approach of bringing deep learning to the life sciences is CDeep3M<sup>8</sup> which ships a set of command line tools and tutorials for training and applying a residual inception network architecture for 3D image segmentation. They specifically address researchers with occasional need for deep learning by providing a cloud-based setup that does not require a local GPU.

The present work provides a generic deep-learning-based software package for cell detection and cell segmentation. For our successful U-Net<sup>3</sup> encoder-decoder network architecture, which has already achieved top ranks in biomedical data analysis benchmarks<sup>9</sup> and which has been the basis of many deep learning models in biomedical image analysis, we developed an interface that runs as plugin in the popular ImageJ software<sup>10</sup> (Supplementary note 1, Supplementary Software 1–4). In contrast to all previous software packages of such kind, our U-Net can be trained and adapted to new data and tasks by the users themselves using the familiar ImageJ interface (Fig. 1). This enables application of U-Net to a wide range of tasks and makes it accessible for a wide set of researchers,

---

<sup>a</sup><http://ilastik.org>

<sup>b</sup>[https://imagej.net/Trainable\\_Weka\\_Segmentation](https://imagej.net/Trainable_Weka_Segmentation)

<sup>c</sup><https://www.drvtechnologies.com/aivia6>

<sup>d</sup><http://cellprofiler.org/>

<sup>e</sup><https://imagej.net/ImgLib2>

who do not have experience with deep learning. For more experienced users the plugin offers an interface to adapt aspects of the network architecture, and to train on data sets from completely different domains. The software comes with a step-by-step protocol and tutorial that shows how to annotate the data for adapting the network and that indicates typical pitfalls (Supplementary Note 2).

U-Net applies to general pixel classification tasks in flat images or volumetric image stacks with one or multiple channels. Such tasks include detection and counting of cells, i.e., prediction of a single reference point per cell, and segmentation, i.e., delineation of the outline of individual cells. These tasks are a superset of the more wide-spread classification tasks, where the object of interest is already localized and only its class label must be inferred. Although adaptation to the detection and segmentation of arbitrary structures in biological tissue is possible given corresponding training data, our experiments focused on cell images, where the flexibility of U-Net is shown on a wide set of common quantification tasks, including detection in multi-channel fluorescence data of dense tissue, and segmentation of cells recorded with different imaging modalities in 2D and 3D (Fig. 2 and Supplementary Note 3). In all shown cases, the U-Net would have saved the responsible scientists much work. As cross-modality experiments show, the diversity in biological samples is too large to ship a single software that can cover them all out-of-the-box (Supplementary note 3.2.2). Thanks to the learning approach, U-Net's applicability is extended from a set of special cases to an unlimited number of experimental settings. The exceptionally precise segmentation results on volumetric brightfield data, the annotation of which gets even human experts to their limits, is a particular strong demonstration of the capabilities of a deep-learning-based automated quantification software (Fig. 2d and Supplementary Note 3.2.4).

Critics often stress the need for huge amounts of training data to train a deep network. In case of U-Net, a base network trained on a diverse set of data and our special data augmentation strategy allow adaptation to new tasks with only one or two annotated images (Online Methods). Only in particularly hard cases, more than 10 training images are needed. Judgment whether a model is adequately trained requires evaluation on a held-out validation set and successive addition of training data until no significant improvement on the validation set can be observed (Supplementary Note 3.2.3). If computational resources allow, cross validation with randomly allocated train/test splits avoids selection of a non-representative validation set at the expense of multiple trainings. The manual effort that must be invested to adapt U-Net to a task at hand is typically much smaller than what is needed for a minimal statistical analysis of the experimental data. In addition, U-Net offers the possibility to run on much larger sample sizes without causing any additional effort. This makes it particularly well-suited for automated large scale sample preparation and recording setups, which are likely to become more and more common in the coming years.

U-Net was optimized for usability in the life sciences. The integration of the software in ImageJ and a step-by-step tutorial make deep learning available to scientists without a computer science background (Supplementary Videos 1–4). Importantly, the computational load is practicable for common life science laboratory settings (Supplementary Note 1). Adaptation of the network to new cell types or imaging modalities ranges from minutes to a few hours on a single computer with a consumer GPU. If the dedicated compute hardware is not available in the lab, common cloud services can be used.

Our experiments provide evidence that U-Net yields results comparable in quality to manual annotation. One special feature of U-Net compared to other automatic annotation tools is the individual

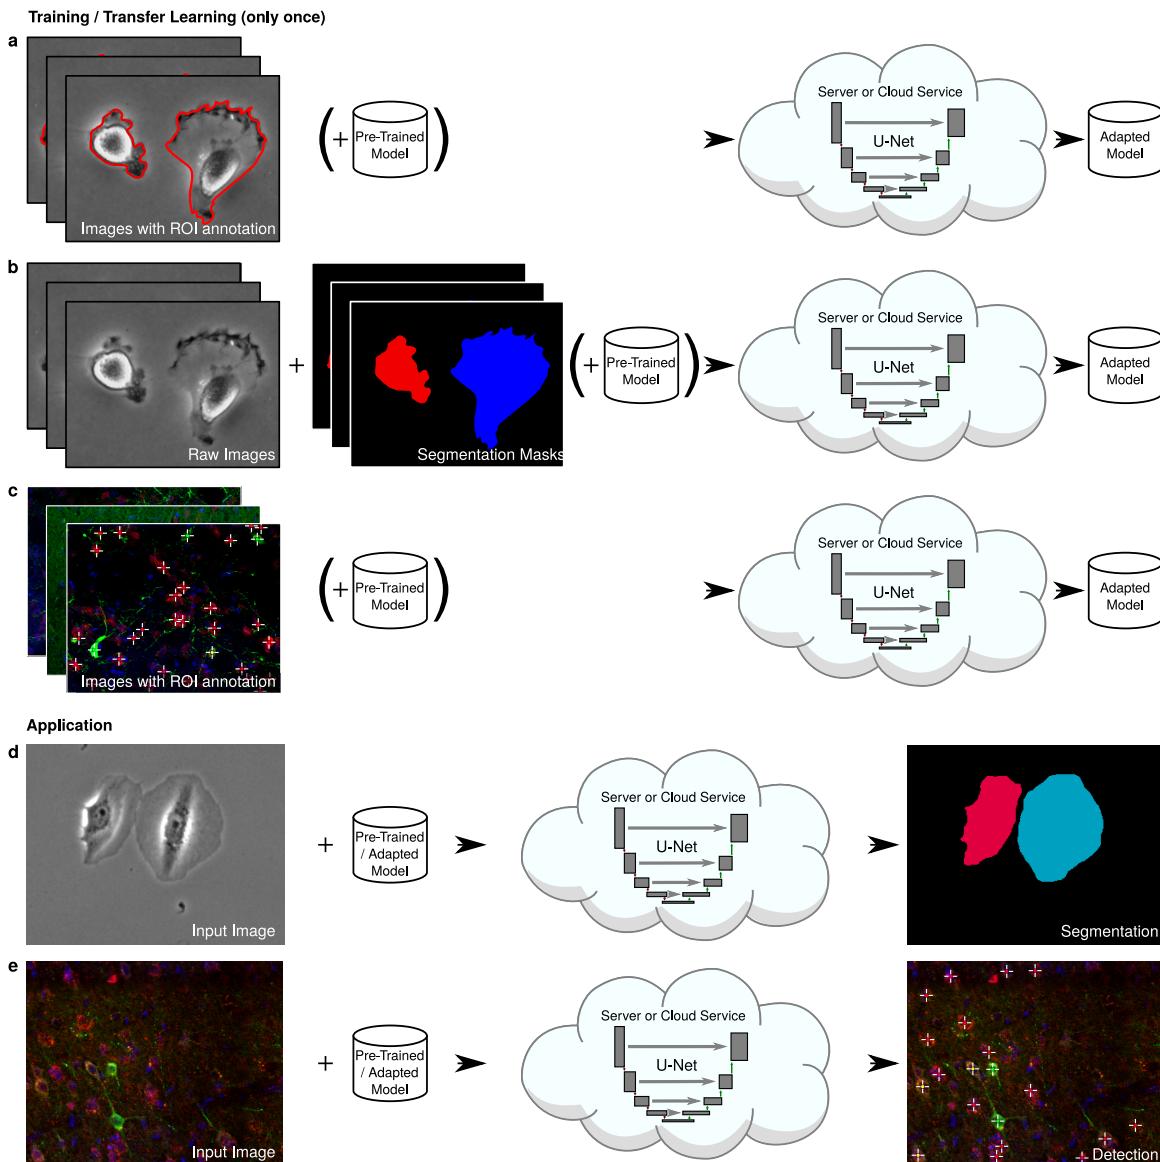


Figure 1: Pipelines of the U-Net software. Left to right: input images and model → network training/application on the local machine, a dedicated remote server or a cloud service → generated output. **(a–c)** Adaptation of the U-Net to newly annotated data using transfer learning; **(a)** Segmentation with regions of interest (ROI) annotations; **(b)** Segmentation with segmentation mask annotations. **(c)** Detection with multi-point annotations; **(d–e)** Application of the pre-trained or adapted U-Net to newly recorded data; **(d)** Segmentation; **(e)** Detection.

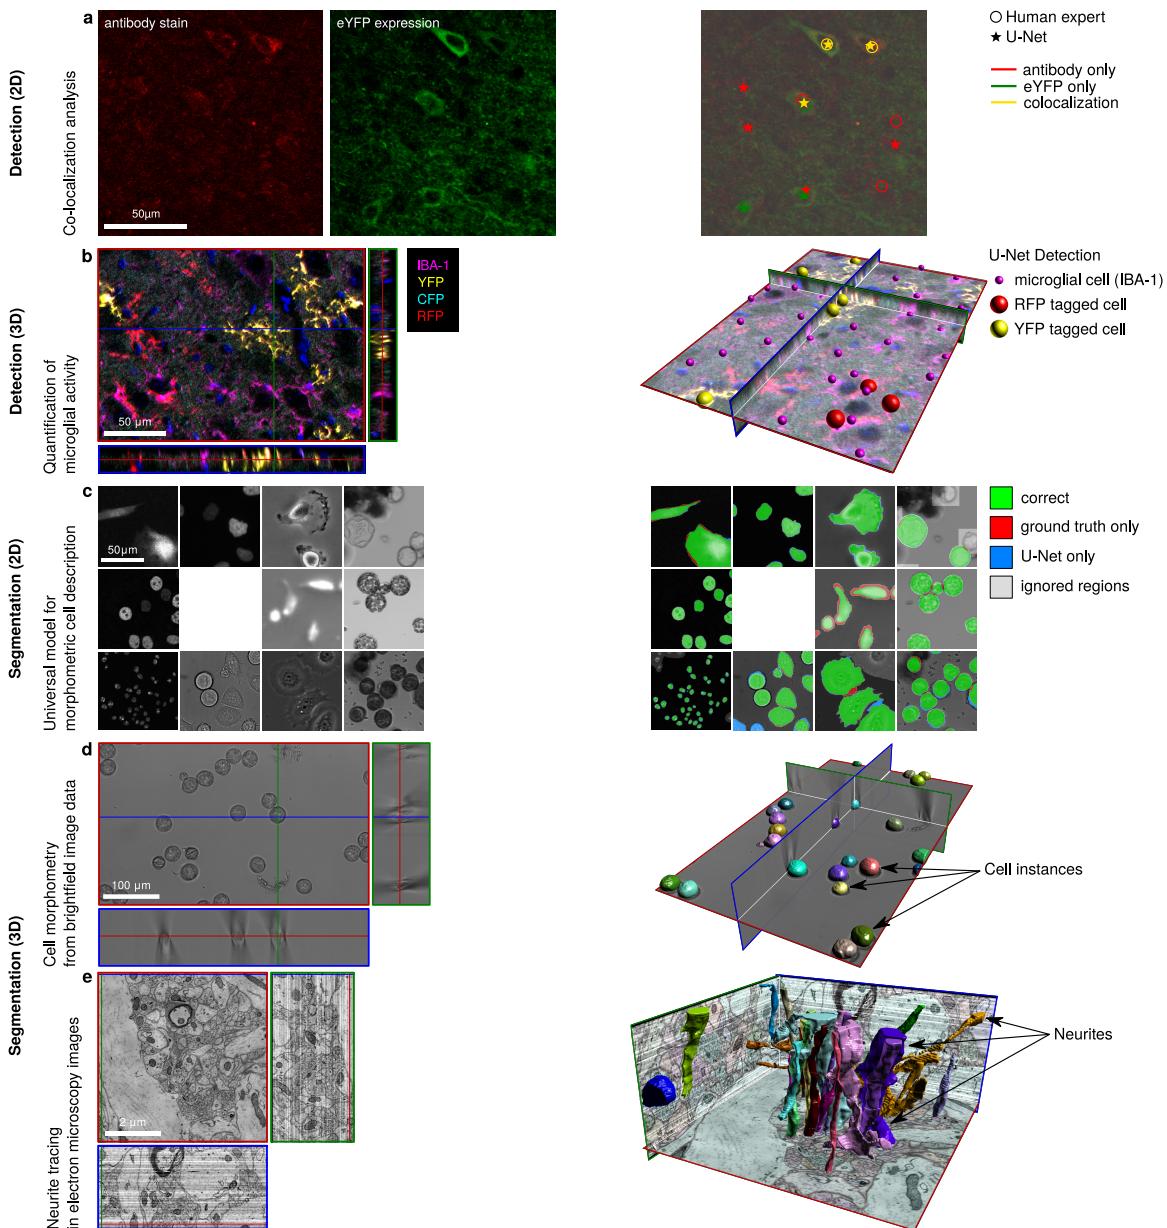


Figure 2: Example applications of the U-Net for 2D and 3D detection and segmentation. **left:** raw data; **right:** U-Net output (with comparison to human annotation in the 2D cases). **(a)** Detection of co-localization in two-channel epi-fluorescence images; **(b)** Detection of xFP-tagged microglial cells in five-channel confocal image stacks. Magenta: all microglia; red, green, cyan: confetti markers; blue: nucleus stain. **(c)** Cell segmentation in 2D images from fluorescence, DIC, phase contrast and brightfield microscopy using one joint model. **(d)** Cell segmentation in 3D brightfield image stacks. **(e)** Neurite segmentation in electron microscopy stacks.

influence of the annotator. This is advantageous as typically researchers develop individual protocols in which multiple parameters are taken into account without being explicitly mentioned. Due to the complexity of these labeling rules they cannot be reproduced by common automatic labeling tools. This advantage can also be turned into a disadvantage: U-Net learns from the provided examples. If these examples are not representative for the actual task or if the manual annotation in these examples is of low quality and inconsistent, U-Net will either fail to train or will reproduce inconsistent annotations on new data. This could also serve as a quality check of the manual annotations. Together, U-Net cannot correct for low-quality human annotations but is a tool to apply individual labeling rules to large data sets and thereby can save manual annotation effort in a vast variety of quantification tasks.

## 1 Acknowledgements

This work was supported by the German Federal Ministry for Education and Research (BMBF) through the MICROSYSTEMS project (0316185B): TF, AD; and the Bernstein Award 2012 (01GQ2301): ID; by the Federal Ministry for Economic Affairs and Energy (ZF4184101CR5): AB; by the Deutsche Forschungsgemeinschaft (DFG) through the collaborative research center KIDGEM (SFB 1140): DM, ÖC, TF, OR; the clusters of excellence BIOSS (EXC 294): TF, DM, RB, AA, YM, DS, TLT, MP, KP, MS, TB, OR; and BrainLinks-Brain-Tools (EXC 1086): ZJ, KS, ID, TB; as well as grants DI 1908/3-1: JD; DI 1908/6-1: ZJ, KS; and DI 1908/7-1: ID; by the Swiss National Science Foundation (SNF Grant 173880): AA; by the ERC Starting Grant OptoMotorPath (338041): ID; and by the FENS-Kavli Network of Excellence (FKNE): ID. We thank F. Prósper, E. Bártová, V. Ulman, D. Svoboda, G. van Cappellen, S. Kumar, T. Becker and the Mitocheck consortium for providing a rich diversity of datasets through the ISBI segmentation challenge. We thank P. Fischer for manual image annotations. We thank S. Wrobel for tobacco microspore preparation.

## 2 Author Contributions

TF, DM, RB, YM, ÖC, TB and OR selected and designed the computational experiments.

TF, RB, DM, YM, AB and ÖC performed the experiments, RB, DM, YM, AB (2D) and TF, ÖC (3D).

R.B., Ö.C., A.A., T.F. and O.R. implemented the U-Net extensions into caffe.

T.F. designed and implemented the Fiji plugin.

D.S. and M.S. selected, prepared and recorded the Keratinocyte dataset PC3-HKPV.

T.F. and O.R. prepared the airborne pollen dataset BF1-POL.

A.D., S.W., O.T., C.D.B. and K.P. selected, prepared and recorded the protoplast and microspore datasets BF2-PPL and BF3-MiSp.

T.L.T. and M.P. prepared, recorded and annotated the data for the microglial proliferation experiment.

J.D. and Z.J. selected, prepared, and recorded the optogenetic data set.

I.D., J.D., and Z.J. manually annotated the optogenetic data set.

I.D., T.F., D.M., R.B., Ö.C., T.B. and O.R. wrote the manuscript.

### 3 Competing financial interests

We declare no competing financial interests.

### 4 Ethics compliance statement

We declare that we complied with all relevant ethical regulations.

## References

1. Sommer, C., Straehle, C., Koethe, U. & Hamprecht, F. A. *Ilastik: Interactive learning and segmentation toolkit in Biomedical Imaging: From Nano to Macro, 2011 IEEE International Symposium on Biomedical Imaging* (2011), 230–233.
2. Arganda-Carreras, I. *et al.* Trainable Weka Segmentation: a machine learning tool for microscopy pixel classification. *Bioinformatics* **33**, 2424–2426 (2017).
3. Ronneberger, O., Fischer, P. & Brox, T. *U-Net: Convolutional Networks for Biomedical Image Segmentation in Medical Image Computing and Computer-Assisted Intervention (MICCAI) 9351* (2015), 234–241.
4. Rusk, N. Deep learning. *Nature Methods* **13**, 35–35 (2016).
5. Webb, S. Deep Learning For Biology. *Nature* **554**, 555–557 (2018).
6. Sadanandan, S. K., Ranefall, P., Le Guyader, S. & Carolina, W. Automated Training of Deep Convolutional Neural Networks for Cell Segmentation. *Scientific Reports* **7**, 7860 (2017).
7. Weigert, M. *et al.* Content-Aware Image Restoration: Pushing the Limits of Fluorescence Microscopy tech. rep. (Center for Systems Biology Dresden (CSBD), Dresden, Germany, 2017).
8. Haberl, M. G. *et al.* CDeep3 – Plug-and-Play cloud-based deep learning for image segmentation. *Nature Methods* **15**, 677–680 (2018).
9. Ulman, V. *et al.* An objective comparison of cell-tracking algorithms. *Nature Methods* **14**, 1141–1152 (2017).
10. Schneider, C. A., Rasband, W. S. & Eliceiri, K. W. NIH Image to ImageJ: 25 years of image analysis. *Nature Methods* **9**, 671–675 (2012).
11. Çiçek, Ö., Abdulkadir, A., Lienkamp, S. S., Brox, T. & Ronneberger, O. *3D U-Net: Learning dense volumetric segmentation from sparse annotation in Medical Image Computing and Computer-Assisted Intervention (MICCAI) 9901* (2016), 424–432.
12. He, K., Zhang, X., Ren, S. & Sun, J. Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification. *CoRR* **abs/1502.01852** (2015).
13. Everingham, M., Van Gool, L., Williams, C. K. I., Winn, J. & Zisserman, A. The Pascal Visual Object Classes (VOC) Challenge. *International Journal of Computer Vision* **88**, 303–338 (2010).
14. Maška, M. *et al.* A benchmark for comparison of cell tracking algorithms. *Bioinformatics* **30**, 1609–1617 (2014).

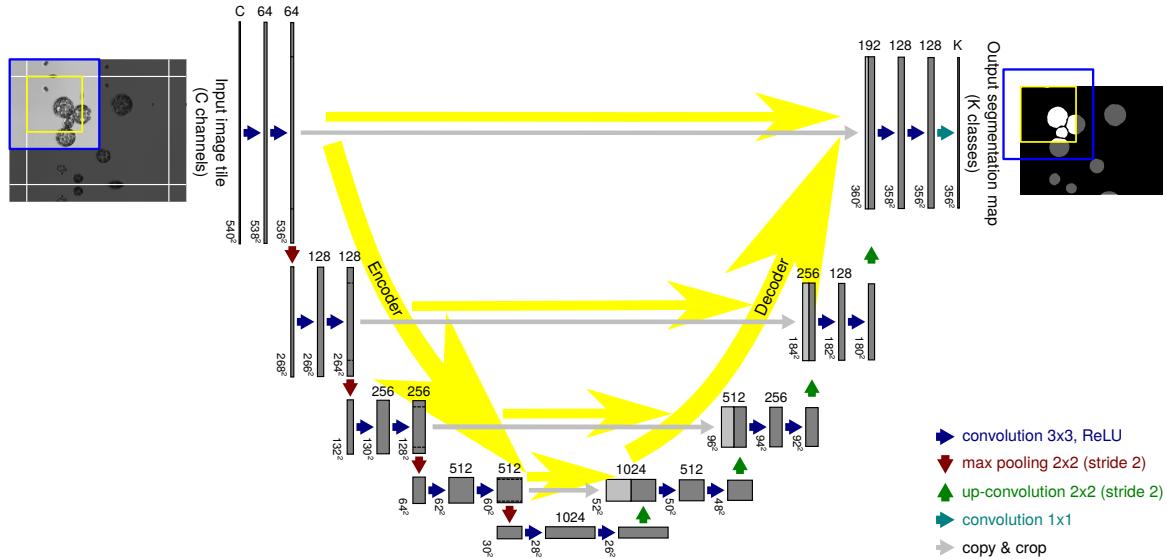


Figure 3: The U-Net architecture at the example of the 2D cell segmentation network. An image tile with  $C$  channels is input to the network on the left (blue box). The  $K$ -class soft-max segmentation is output on the right (yellow box). Blocks show the computed feature hierarchy. Numbers atop each network block: number of feature channels; numbers left to each block: spatial feature map shape in pixels. Yellow arrows: Data flow.

## Online Methods

### M1 Network architecture

The U-Net is an encoder-decoder-style neural network solving semantic segmentation tasks end-to-end (Fig. 3). Its 26 hidden layers can be logically grouped in two parts: An encoder that takes an image tile as input and successively computes feature maps at multiple scales and abstraction levels yielding a multi-level, multi-resolution feature representation; a decoder that takes the feature representation and classifies all pixels/voxels at original image resolution in parallel. Layers in the decoder gradually synthesize the segmentation starting at low-resolution feature maps (describing large scale structures) up to full-resolution feature maps (describing fine scale structures).

The encoder is a typical convolutional neural network. It consists of the repeated application of two convolutions (without padding), each followed by a leaky rectified linear unit (ReLU) with leakage factor 0.1 and a max-pooling operation with stride two halving the resolution of the resulting feature map. Convolutions directly following down-sampling steps double the number of feature channels.

The decoder consists of repeated application of an up-convolution (an up-convolution is equivalent to a bed-of-nails up-sampling by a factor of two followed by a convolution with edge length two) that halves the number of feature channels, then concatenation with the cropped encoder feature map at corresponding resolution and two convolutions each followed by leaky ReLUs with leakage factor 0.1.

At the final layer a  $1 \times 1$  convolution is used to map feature vectors to the desired number of classes

$K$ .

For training, the final  $K$ -class scores are transformed to pseudo-probabilities with soft-max before they are compared to the ground truth annotations in one-hot encoding using cross entropy.

If not explicitly mentioned, operations are isotropic and convolutions use a kernel edge length of three pixels. We only use the valid part of each convolution, *i.e.* the segmentation map only contains the pixels, for which the full context is available in the input image tile.

This fully convolutional network architecture allows to freely choose the input tile shape with the restriction that edge lengths of feature maps prior to max pooling operations must be even. This implies the possibility to process arbitrarily large images using an overlap-tile strategy where the stride is given by the output tile shape. To predict pixels in the border region of an image, missing context is extrapolated by mirroring the input image.

The ImageJ plugin presents only valid tile shapes to the user. For 2D networks the user can also let the plugin automatically choose an optimal tiling based on the available GPU memory. For this the amount of memory used by the given architecture for different tile sizes must be measured on the client PC beforehand and corresponding information stored to the model definition file using the provided MATLAB script `caffe-unet/matlab/unet/measureGPUMem.m`.

## M1.1 3D U-Net

The 3D U-Net architecture for volumetric inputs is only a slight modification of its 2D counterpart. First, input tiles have a shape of  $236 \times 236 \times 100$  voxels (vx). Secondly, due to memory limitations, we half the number of output features of each convolution layer except for the last up-convolution step that only produces 32 channels. Thirdly, to deal with anisotropic image voxel sizes, convolutions and pooling at the finest resolution levels are 2D until voxel extents are roughly isotropic. At all remaining resolution levels operations are applied along all three dimensions.

## M2 Weighted soft-max cross-entropy loss

We use pixel-weighted soft-max cross-entropy loss to allow to change the influence of imbalanced classes in semantic segmentation. The loss is computed as

$$l(I) := - \sum_{\mathbf{x} \in \Omega} w(\mathbf{x}) \log \frac{e^{\hat{y}_{y(\mathbf{x})}(\mathbf{x})}}{\sum_{c=0}^C e^{\hat{y}_c(\mathbf{x})}}$$

where  $\mathbf{x}$  is a pixel in image domain  $\Omega$ ,  $\hat{y}_c : \Omega \rightarrow \mathbb{R}$  is the predicted score for class  $c \in \{0, \dots, C\}$ ,  $C$  is the number of classes,  $y : \Omega \rightarrow \{0, \dots, C\}$  is the true class label. With this,  $\hat{y}_{y(\mathbf{x})}(\mathbf{x})$  is the predicted score for the ground truth class at position  $\mathbf{x}$ . As defined above,  $w : \Omega \rightarrow \mathbb{R}_{\geq 0}$  is the pixel-wise loss weight.

We employ loss weighting to optimize class-balancing and handle regions without annotations (in the following termed "ignored regions") using weight map  $w_{\text{bal}}$ . We additionally enforce instance separation using weight  $w_{\text{sep}}$  as described in the following sections. The final loss weights are then given by

$$w(\mathbf{x}) := w'_{\text{bal}} + \lambda w_{\text{sep}}$$

where  $\lambda \in \mathbb{R}_{\geq 0}$  controls the importance of instance separation.

## M2.1 Class balancing and regions with unknown ground truth

In our images most pixels belong to the background class, and in many cases this class is homogeneous and easy to learn. Therefore, we reduce the weight of background pixels by factor  $v_{\text{bal}} \in [0, 1]$  compared to foreground pixels resulting in the class balancing weight function

$$w_{\text{bal}}(\mathbf{x}) := \begin{cases} 1 & y(\mathbf{x}) > 0 \\ v_{\text{bal}} & y(\mathbf{x}) = 0 \\ 0 & y(\mathbf{x}) \text{ unknown (i.e. ignored regions).} \end{cases}$$

We observed slightly better segmentation results when replacing the step-shaped cutoff at the edges of foreground objects by a smoothly decreasing Gaussian function for the weighted loss computation, therefore we define

$$w'_{\text{bal}}(\mathbf{x}) := \begin{cases} 1 & y(\mathbf{x}) > 0 \\ v_{\text{bal}} + (1 - v_{\text{bal}}) \cdot \exp\left(-\frac{d_1^2(\mathbf{x})}{2\sigma_{\text{bal}}^2}\right) & y(\mathbf{x}) = 0 \\ 0 & y(\mathbf{x}) \text{ unknown (i.e. ignored)} \end{cases}$$

where  $d_1(\mathbf{x})$  is the distance to the closest foreground object and  $\sigma_{\text{bal}}$  is the standard deviation of the Gaussian.

For the **detection** tasks we added rings/spherical shells of ignored regions with outer radius  $r_{\text{ign}}$  around annotated locations to both  $w_{\text{bal}}$  and  $w'_{\text{bal}}$ . These regions help to stabilize the detection process in case of inaccurate annotations, but are not essential to the approach.

## M2.2 Instance segmentation

Semantic segmentation classifies pixels/voxels of an image. Therefore, touching objects of the same class will end up in one joint segment. Most often we want to measure properties of individual object instances (*e.g.* cells). Therefore we must turn the pure semantic segmentation into an instance-aware semantic segmentation. For this we insert an artificial one pixel wide background ridge between touching instances in the ground truth segmentation mask (Fig. 4a).

To force the network to learn these ridges we increase their weight in the loss computation such that the thinnest ridges get the highest weights. We approximate the ridge width at each pixel by the sum of the distance  $d_1$  to its nearest instance and the distance  $d_2$  to its second-nearest instance. From this we compute the weight map as

$$w_{\text{sep}}(\mathbf{x}) := \exp\left(-\frac{(d_1(\mathbf{x}) + d_2(\mathbf{x}))^2}{2\sigma_{\text{sep}}^2}\right),$$

which decreases following a Gaussian curve with standard deviation  $\sigma_{\text{sep}}$  (Fig. 4b).

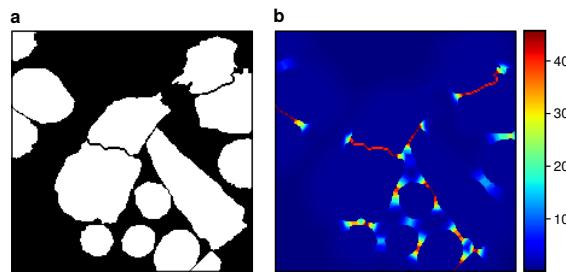


Figure 4: Separation of touching cells using pixel-wise loss weights. **(a)** Generated segmentation mask with one pixel wide background ridge between touching cells (white: foreground, black: background) **(b)** Map showing pixel-wise loss weights to enforce the network to separate touching cells.

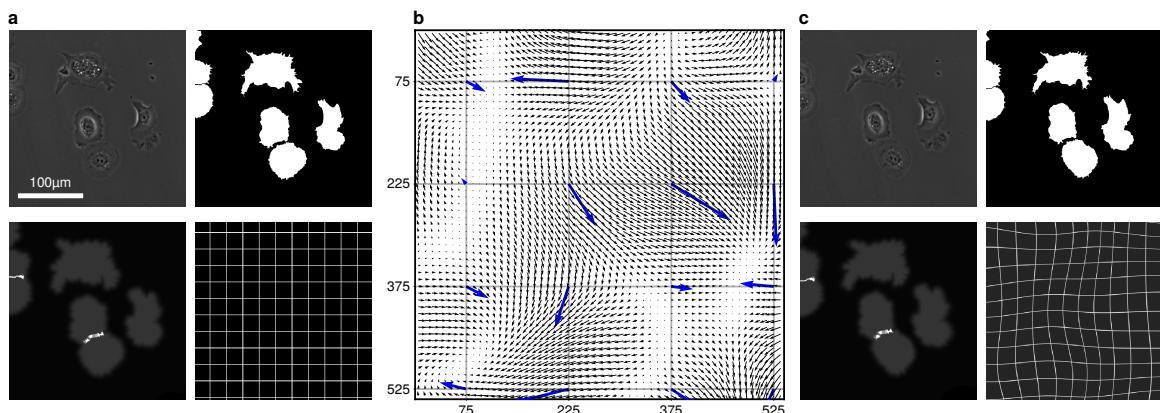


Figure 5: Training data augmentation by random smooth elastic deformation. **(a)** Upper left: Raw image; Upper right: Labels; Lower Left: Loss Weights; Lower Right: 20 $\mu\text{m}$  grid (for illustration purpose only) **(b)** Deformation field (black arrows) generated using bicubic interpolation from a coarse grid of displacement vectors (blue arrows; magnification: 5 $\times$ ). Vector components are drawn from a Gaussian distribution ( $\sigma = 10\text{px}$ ). **(c)** Back-warp transformed images of (a) using the deformation field.

### M3 Tile sampling and augmentation

Data augmentation is essential to teach the expected appearance variation to the neural network with only few annotated images. To become robust to translations and to focus on relevant regions in the images, we first randomly draw spatial locations of image tiles that are presented to the network from a user-defined probability density function (pdf). For all our experiments we used the normalized weight map  $w_{\text{bal}}$  for sampling the spatial location, *i.e.* foreground objects are presented to the network ten times more often (according to our selection of  $v_{\text{bal}} = 0.1$ ) than background regions and tiles centered around an ignore region are never selected during training. Then, we draw a random rotation angle (around the optical axis in 3D) from a uniform distribution in a user-defined range. Finally, we generate a smooth deformation field by placing random displacement vectors with

user-defined standard deviation of the magnitude for each component on a very coarse grid. These displacement vectors are used to generate a dense full-resolution deformation field using bicubic interpolation (Fig. 5). Rigid transformations and elastic deformations are concatenated to look-up intensities in the original image during tile generation.

We additionally apply a smooth strictly increasing intensity transformation to become robust to brightness and contrast changes. The intensity mapping curve is generated from a user-defined number of equidistant control points in the normalized [0, 1] source intensity range. Target intensities at the control points are drawn from uniform distributions with user defined ranges. The sampling process enforces intensities at subsequent control points to increase and spline interpolation between the control points ensures smoothness.

All data augmentation is applied on-the-fly to input image, labels and weight maps during network training<sup>11</sup>.

## M4 Training

All networks were trained on an nVidia TITAN X with 12GB GDDR5 RAM using cuda 8 and cuDNN 6 with caffe after applying our proposed extensions. The initial network parameters are drawn from a Gaussian distribution with standard deviation  $\sigma = \sqrt{2/n_{in}}$ , where  $n_{in}$  is the number of inputs of one neuron of the respective layer<sup>12</sup>.

For all experiments, raw image intensities per channel were normalized to the [0, 1] range before training using

$$\hat{I} := \frac{I - \min\{I\}}{\max\{I\} - \min\{I\}}$$

where  $I$  is the raw and  $\hat{I}$  the normalized intensity.

## M5 Transfer learning

Adaptation of a pre-trained U-Net to a new dataset using annotated data is called transfer learning or finetuning. Transfer learning leverages the knowledge about the different cell datasets already learned by the U-Net and usually requires considerably less annotated data and training iterations compared to training from scratch.

Transfer to a new dataset is based on the same training protocol as described above. All the user must provide are raw images and corresponding annotations as ImageJ ROIs or pixel masks (Detection: One Multi-point ROI per class and image; Segmentation: One regional ROI per object or pixel masks) (Supplementary Section B). The plugin performs image re-scaling, pixel weight  $w(\mathbf{x})$  generation, and parametrization of the caffe training software.

## M6 Evaluation Metrics

### M6.1 Object Intersection over Union

The intersection over union (IoU) measures how well a predicted segmentation matches the corresponding ground truth annotation by dividing the intersection of two segments by their union. Let  $\mathcal{S} := \{s_1, \dots, s_N\}$  be a set of  $N$  pixels. Let  $\mathcal{G} \subset \mathcal{S}$  be the set of pixels belonging to a ground truth segment and  $\mathcal{P} \subset \mathcal{S}$  be the set of pixels belonging to the corresponding predicted segment. The IoU is defined as

$$M_{\text{IoU}}(\mathcal{G}, \mathcal{P}) = \frac{|\mathcal{G} \cap \mathcal{P}|}{|\mathcal{G} \cup \mathcal{P}|}.$$

IoU is a widely used measure, *e.g.* in the Pascal VOC challenge<sup>13</sup> or the ISBI cell tracking challenge<sup>14</sup>.  $M_{\text{IoU}} \in [0, 1]$  with 0 meaning no overlap at all and 1 meaning a perfect match. In our experiments, a value of  $\sim .7$  indicates a good segmentation result, and a value of  $\sim .9$  is close to human annotation accuracy.

We first determine the predicted objects by performing a connected component labelling with 8-neighborhood on the binary output of the U-Net. This yields candidate segments  $S_i$ . Then we compute the IoU for every pair of output and ground truth segments  $G_j$ :  $M_{\text{IoU}}(S_i, G_j)$  and apply the Hungarian algorithm on  $M_{\text{IoU}}$  to obtain 1:1 correspondences maximizing the average IoU. Unmatched segments or matches with zero IoU are considered false positives and false negatives, respectively. The average IoU is computed based on the ground truth annotations, *i.e.* false negatives contribute to the average IoU with a value of 0. False positives, however, are not considered in the average IoU.

In the supplementary material we provide a detailed analysis of false positive segmentations produced by the U-Net.

### M6.2 F-measure

We measure the quality of a detector using the balanced  $F_1$  score (denoted as F-measure) which is the harmonic mean of precision and recall. Let  $G$  be the set of true object locations and  $P$  be the set of predicted object locations. Then precision and recall are defined as

$$\text{Precision} := \frac{|G \cap P|}{|P|} \quad \text{and} \quad \text{Recall} := \frac{|G \cap P|}{|G|}.$$

The F-measure is then given by

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}.$$

Pixel-accurate object localization is almost impossible to reach, therefore, we introduce a tolerance  $d_{\text{match}}$  for matching detections and annotations. To get one-to-one correspondences, we first compute the pairwise distances of ground truth positions to detections. Predictions with distance greater

than  $d_{\text{match}}$  to any ground truth position can be directly classified as false positive. Similarly ground truth positions without detections in the  $d_{\text{match}}$  range are classified as false negatives. We apply the Hungarian algorithm to the remaining points to obtain an assignment of predictions to ground truth positions minimizing the total sum of distances of matched positions. Correspondences with distance greater than  $d_{\text{match}}$  and unmatched positions in either ground truth or prediction are treated as false negative/false positive detections as before.

## M7 Code availability

We provide pre-built binary versions of the U-Net caffe extensions for Ubuntu Linux 16.04 at <https://lmb.informatik.uni-freiburg.de/resources/opensource/unet/>. We additionally provide our changes to the source code of the publicly available caffe deep learning framework<sup>f</sup> as patch file with detailed instructions on how to apply the patch and build our caffe variant from source in the supplementary material.

Binary installation only requires to unpack the archive and install required third-party libraries which can be done within few minutes on an Ubuntu 16.04 machine depending on your internet connection for fetching the packages.

Building from scratch requires to install the dependent development libraries and checkout the given tagged version of the BVLC caffe master branch and apply the patch. With "normal" internet connection, package installation is a matter of few minutes. Cloning the BVLC master repository requires less than a minute, applying the patch imposes no measurable overhead. Configuring and building the package requires approximately ten to fifteen minutes.

The U-Net segmentation plugin for Fiji/ImageJ is available at <http://sites.imagej.net/Falk/plugins/> or through the ImageJ updater within Fiji. Source code is included in the plugin jar File Unet\_Segmentation.jar. Installation using the Fiji updater requires only a few seconds.

The trained caffe-models for the 2d- and 3d-U-Net are available at <https://lmb.informatik.uni-freiburg.de/resources/opensource/unet/>.

## M8 Data availability

The datasets F1-MSC, F2-GOWT1, F3-SIM, F4-HeLa, DIC1-HeLa, PC1-U373, and PC2-PSC are from the ISBI Cell Tracking Challenge 2015<sup>14</sup>. Information on how to obtain the data can be found at <http://celltrackingchallenge.net/datasets.html> and currently requires free-of-charge registration for the challenge.

The datasets PC3-HKPV, BF1-POL, BF2-PPL, and BF3-MiSp are custom and are available from the corresponding author upon reasonable request.

Datasets for the detection experiments partially contain unpublished sample preparation protocols, and are currently not freely available. Upon protocol publication datasets will be made available on request-basis.

---

<sup>f</sup><https://github.com/BVLC/caffe>